
IERG 4180
Network Software Design and Programming

Client-Server Programming
Part I – Fundamentals

Rev. 1.1 – 28 Sep 2016

Contents

Copyright Jack Y. B. Lee
All Rights Reserved

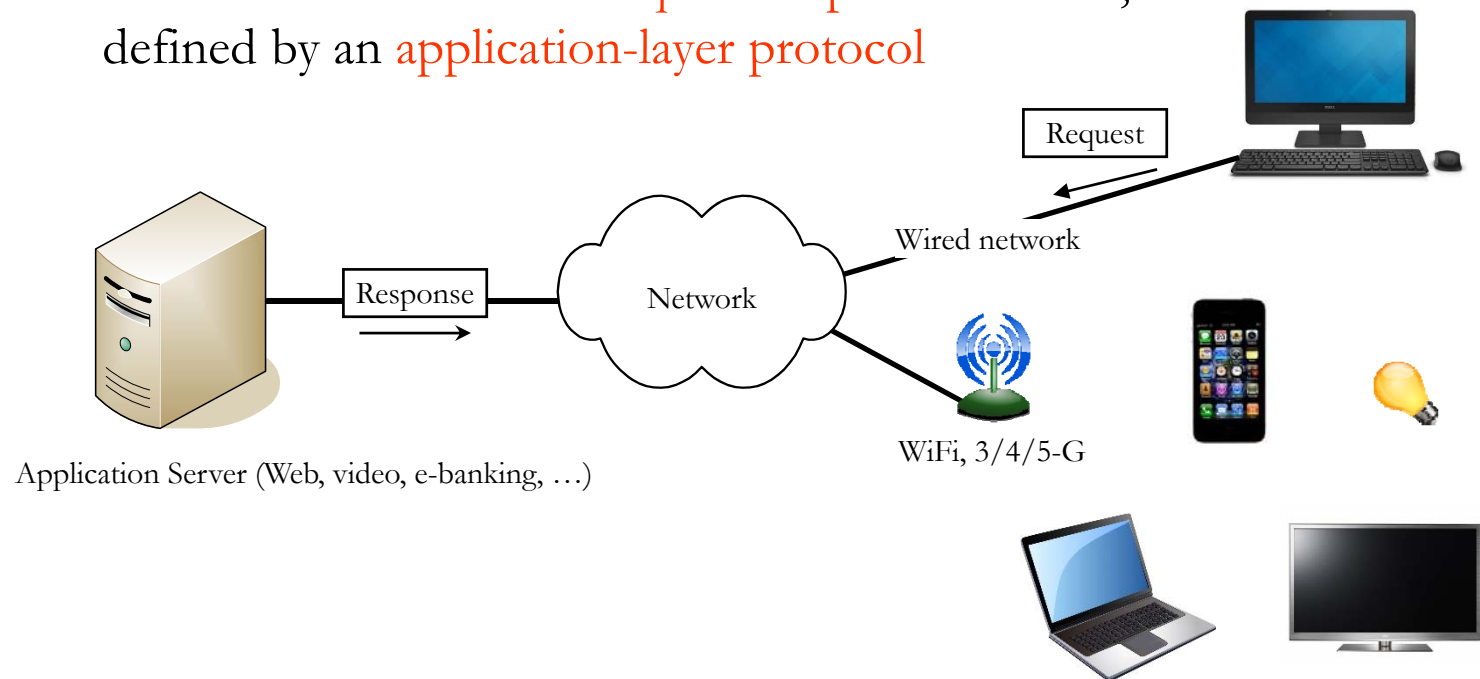
1. Client-Server Model
2. Iterative Server
3. Concurrent Server
4. I/O Multiplexing
5. Non-Blocking I/O with Polling
6. I/O Multiplexing Using `select ()`
7. Concurrent `NetConsole`
8. Final Remarks on `select ()`

1. Client-Server Model

Copyright Jack Y. B. Lee
All Rights Reserved

- Architecture

- ◆ Server **hardware** and **operating system** platform choices are more flexible
- ◆ Interaction is often in a **request-response** manner, defined by an **application-layer protocol**



Clients may run in a wide variety of platforms and connect via different kinds of networks.

1. Client-Server Model

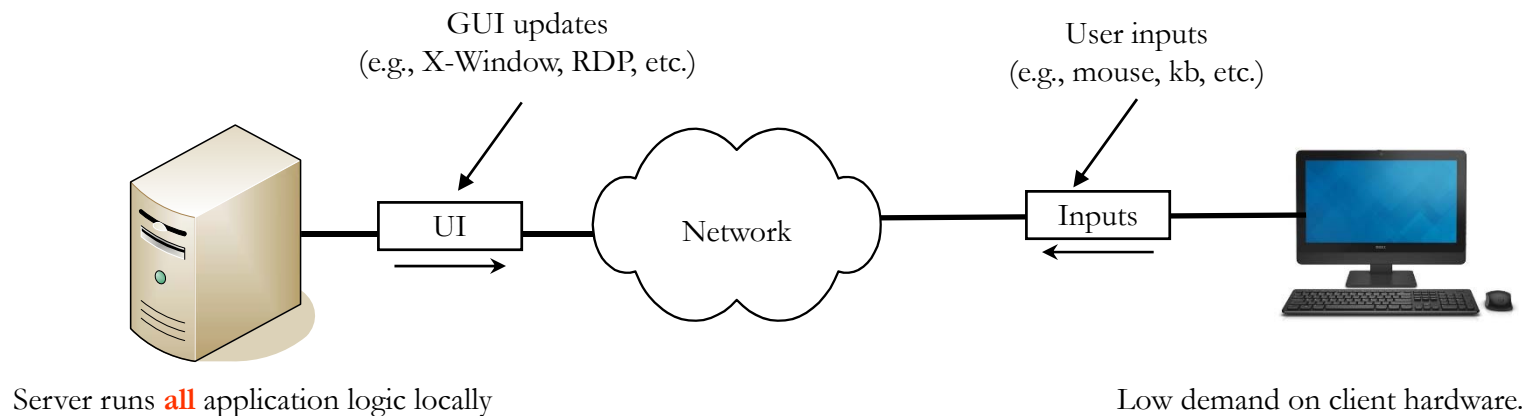
Copyright Jack Y. B. Lee
All Rights Reserved

- Application Server
 - ◆ Able to process requests from **multiple clients simultaneously**
 - ◆ Automatically runs upon operating system startup
 - ◆ Operates continuously as a **long-running process**
 - ◆ Listen for incoming connections/requests using a well-known transport address (e.g., `http://www.cuhk.edu.hk`)
 - ◆ Must fail in a **predictable** manner and can be restarted into a **consistent** state
 - ◆ Public servers are open to network **intruders** and **attacks**
 - ◆ A plan for **capacity scalability** cannot be an after-thought
- Application Client
 - ◆ **Initiates** requests to servers to request data/service
 - ◆ May perform **local processing** independent of server
 - ◆ May **not** require constant connection to server to execute

1. Client-Server Model

Copyright Jack Y. B. Lee
All Rights Reserved

- Contrast to Thin-Client (or Terminal) Model
 - ◆ Client responsible for **user interface rendering and interaction** only
 - ◆ **No transfer** of actual data to client for processing



Examples:

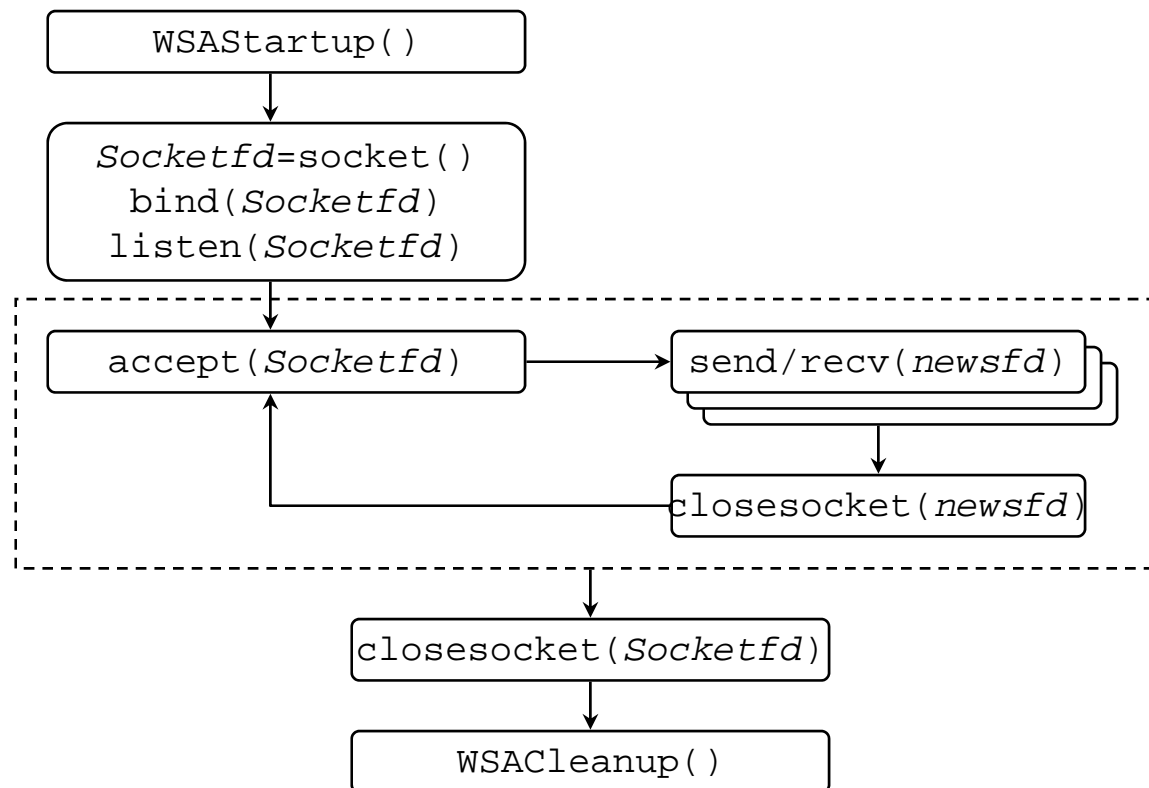
True Client-Server: Playing a multiplayer car racing game using PC connected to a game server.

Terminal Model: Playing games via platforms such as OnLive (acquired by Sony in 2015).

2. Iterative Server

Copyright Jack Y. B. Lee
All Rights Reserved

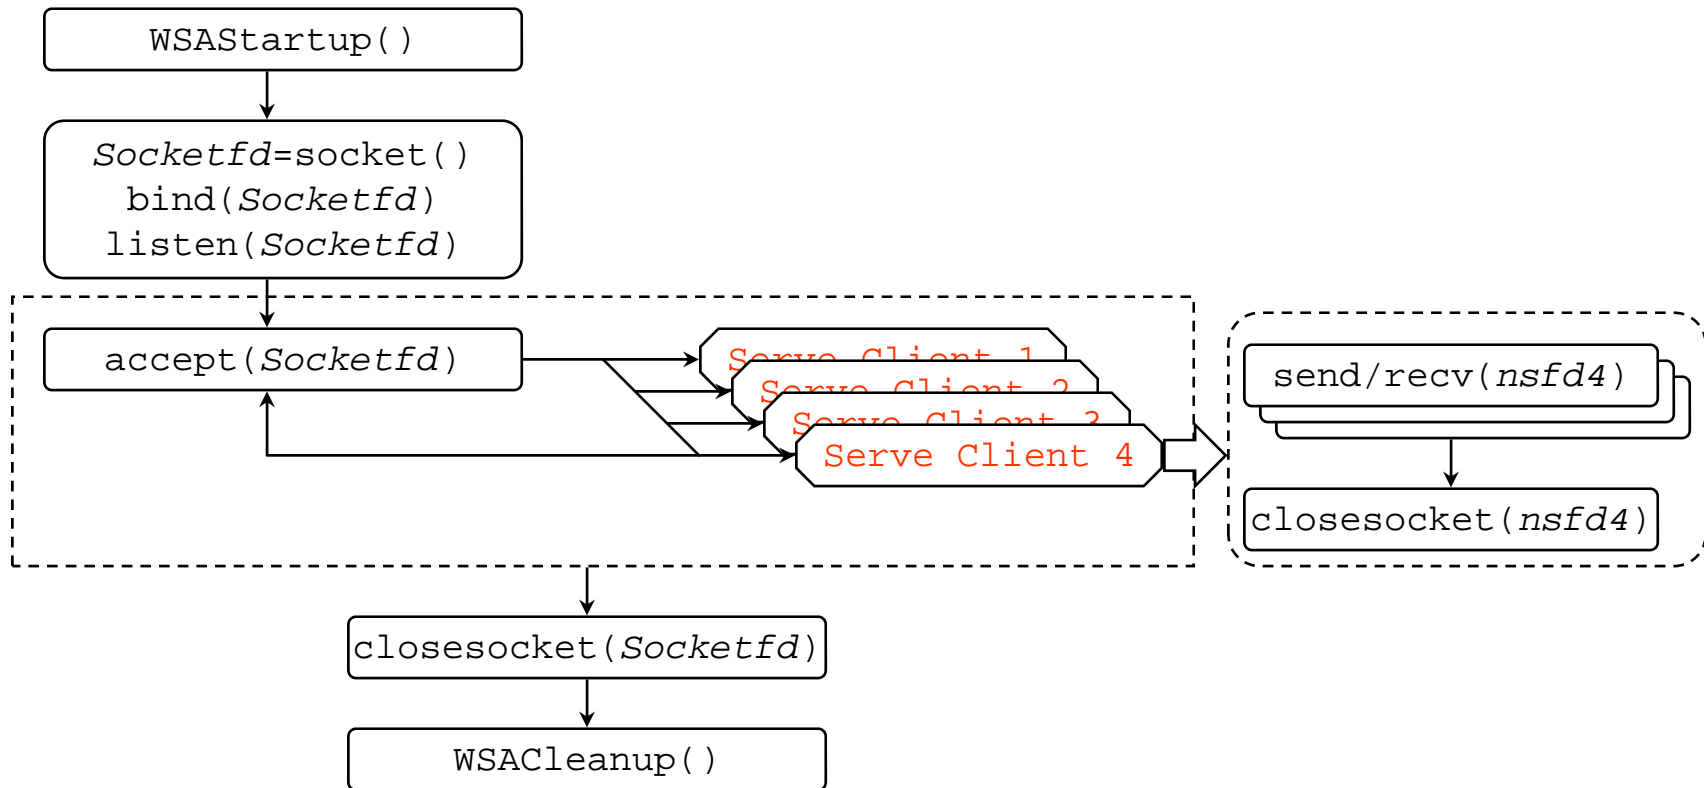
- Serves **one** client at a time (e.g., NetConsole Listener).
 - ◆ Subsequent request will need to **wait** until previous one finishes (how?)



3. Concurrent Server

Copyright Jack Y. B. Lee
All Rights Reserved


- Able to serve **multiple clients concurrently**.
 - ◆ Server resources (e.g., bandwidth, memory) are shared across concurrent client sessions



3. Concurrent Server

Copyright Jack Y. B. Lee
All Rights Reserved


- Synchronous programming and blocking I/O model.
 - ◆ Accepting a second connection:



3. Concurrent Server

Copyright Jack Y. B. Lee
All Rights Reserved

- Synchronous programming and blocking I/O model.
 - ◆ Accepting a second connection:



Question: What is the fundamental problem here?

4. I/O Multiplexing

Copyright Jack Y. B. Lee
All Rights Reserved

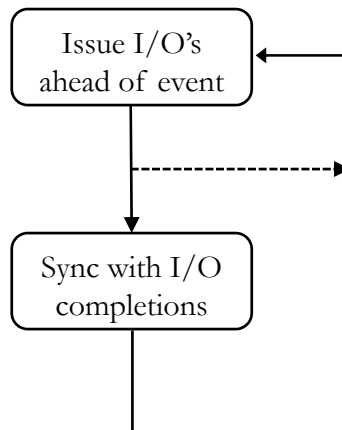
- The Problem
 - ◆ Program execution is driven by process/thread
 - ◆ A process/thread blocked by an I/O call will become **suspended**, and unable to serve other pending I/O calls
- Synchronous Programming Approaches*
 - ◆ Non-blocking I/O with polling (Unix, Windows)
 - ◆ The `select ()` system call (Unix, Windows, also Java)
 - ◆ Multi-Process with blocking I/O (Unix)
 - ◆ Multi-threading with blocking I/O (Unix, Windows, Java)
 - ◆ Overlapped I/O (Windows only)
 - ◆ Combinations of the above

4. I/O Multiplexing

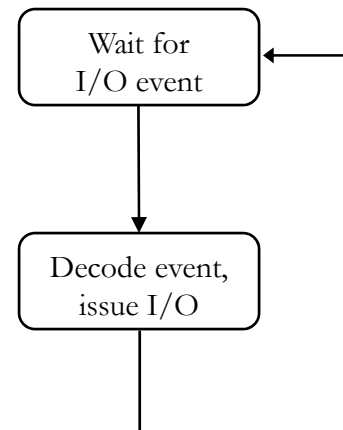
Copyright Jack Y. B. Lee
All Rights Reserved

- Asynchronous Programming Approaches
 - ◆ Message-driven I/O (Windows only)
 - ◆ I/O Completion Port, Registered I/O (Windows only)
 - ◆ Unix Signals SIGIO (Unix only)
 - ◆ Linux `poll` and `epoll`
 - ◆ Combination with multi-process / multi-threading

Synchronous Approaches




Asynchronous Approaches



5. Non-Blocking I/O with Polling

Copyright Jack Y. B. Lee
All Rights Reserved

- Turns a socket into non-blocking mode
 - ◆ On a **per-socket basis**, default is blocking
 - ◆ Operations on non-blocking sockets are guaranteed to not block
 - ◆ Three possible outcomes
 - Completed successfully
 - ◆ e.g., `recv()` returns > 0
 - Failed
 - ◆ e.g., `recv()` returns `SOCKET_ERROR` with `WSAGetLastError() != WSAEWOULDBLOCK`
 - Cannot complete w/o blocking
 - ◆ e.g., `recv()` returns `SOCKET_ERROR` with `WSAGetLastError() == WSAEWOULDBLOCK`



5. Non-Blocking I/O with Polling

Copyright Jack Y. B. Lee
All Rights Reserved

- Turns a socket into non-blocking mode

```
void ConcurrentListenerUsingNonBlockingIOandPolling(char *address, char *port)
{
    /// Step 1: Prepare address structures. ///
    sockaddr_in *TCP_Addr = new sockaddr_in;
    memset (TCP_Addr, 0, sizeof(struct sockaddr_in));
    TCP_Addr->sin_family = AF_INET;
    TCP_Addr->sin_port = htons(atoi(port));
    TCP_Addr->sin_addr.s_addr = INADDR_ANY;

    /// Step 2: Create a socket for incoming connections. ///
    SOCKET Sockfd = socket(AF_INET, SOCK_STREAM, 0);
    bind(Sockfd, (struct sockaddr *)TCP_Addr, sizeof(struct sockaddr_in));
    listen(Sockfd, 5);
    unsigned long ul_val = 1; // 1 means enable
    if (ioctlsocket(Sockfd, FIONBIO, &ul_val) == SOCKET_ERROR) {
        printf("\nioctlsocket() failed. Error code: %i\n", WSAGetLastError());
        return;
    }

    // ...
}
```

Note: For Linux use `ioctl()` instead of `ioctlsocket()`.

5. Non-Blocking I/O with Polling

Copyright Jack Y. B. Lee
All Rights Reserved

- `accept ()` will become non-blocking:

```
/// Step 3: Setup the data structures for multiple connections. ///
const int maxSockets = 10;
SOCKET socketHandles[maxSockets]; // Array for the socket handles
bool    socketValid[maxSockets];  // Bitmask to manage the array
int     numActiveSockets = 1;
for (int i=1; i<maxSockets; i++) socketValid[i] = false;
socketHandles[0] = Sockfd;  socketValid[0] = true;

/// Step 4: Poll all active sockets for data/accept. ///
while (1) {
    // Check for incoming connection first.
    if (socketValid[0]) {
        SOCKET newsfd = accept(Sockfd, 0, 0);
        if (newsfd != SOCKET_ERROR) { // accept() succeeded
            // Append the new entry to the socketHandles[] array//
            socketHandles[numActiveSockets] = newsfd;
            socketValid[numActiveSockets] = true;
            ++numActiveSockets;
            if (numActiveSockets == maxSockets) { // Suspend accept
                socketValid[0] = false;
            }
        } else if (WSAGetLastError() != WSAEWOULDBLOCK) {
            printf("\naccept() failed. Error code: %i\n", WSAGetLastError());
            return;
        } // Note the case for WSAEWOULDBLOCK is ignored.
    }
}
```

5. Non-Blocking I/O with Polling

Copyright Jack Y. B. Lee
All Rights Reserved


- Poll the rest of the active sockets:

```
// ...
for (int i=1; i<numActiveSockets; i++) {
    if (socketValid[i]) {
        // Receive data and display to stdout. //
        char buf[256]; int len=255;
        int numread = recv(socketHandles[i], buf, len, 0);
        if (numread == SOCKET_ERROR) {
            if (WSAGetLastError() != WSAEWOULDBLOCK) {
                printf("\nrecv() failed. Error: %i\n", WSAGetLastError());
                return;
            }
            continue;
        } else if (numread == 0) { // connection closed
            // Pack the socket array
            closesocket(socketHandles[i]);
            --numActiveSockets;
            socketHandles[i] = socketHandles[numActiveSockets];
            socketValid[numActiveSockets] = false;
            if (numActiveSockets == (maxSockets-1))
                socketValid[0] = true;
        } else {
            printf("[%i]: ", i);
            for (int x=0; x<numread; x++) putchar(buf[x]);
        }
    }
} // end-for
```

5. Non-Blocking I/O with Polling

Copyright Jack Y. B. Lee
All Rights Reserved

- Program Flow

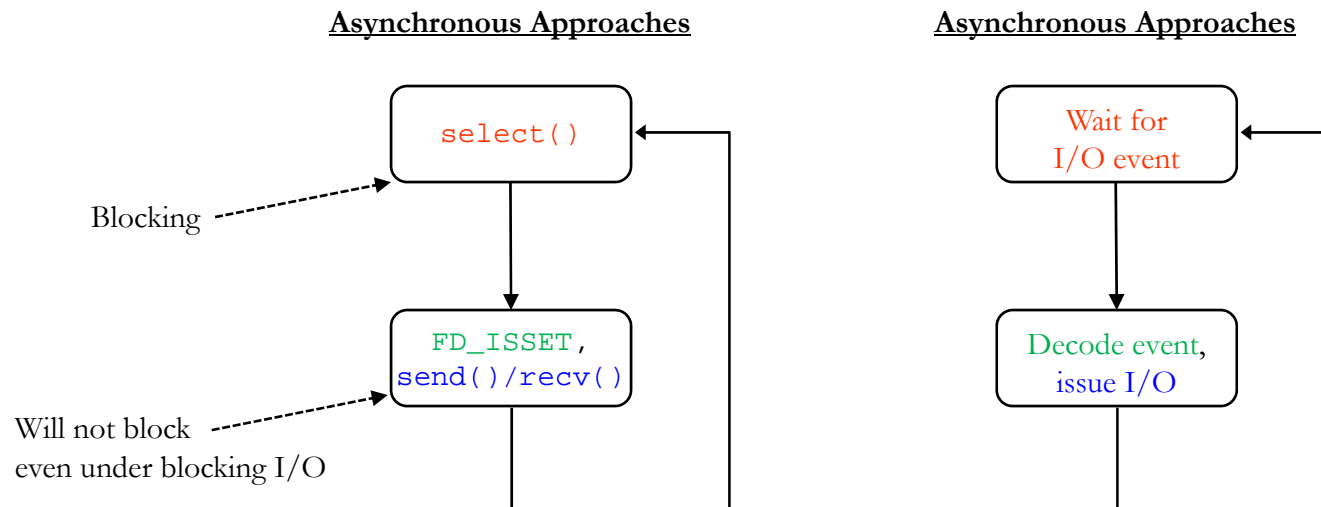


Question: What is the tradeoff of polling non-blocking sockets?

6. I/O Multiplexing Using `select()`

Copyright Jack Y. B. Lee
All Rights Reserved

- Asynchronous Approach
 - ◆ Blocking I/O, single thread



```
int select( int nfds,           // Max fd+1 in Unix/Linux, ignored in windows.
            fd_set* readfds,    // Set of sockets for read events.
            fd_set* writefds,   // Set of sockets for write events.
            fd_set* exceptfds,  // Set of sockets for exceptions/errors.
            const struct timeval* timeout // Max time to block.
);
```

6. I/O Multiplexing Using select ()

Copyright Jack Y. B. Lee
All Rights Reserved

- The select () API function:

```
int select( int nfds,           // Max fd+1 in Unix/Linux, ignored in windows.
            fd_set* readfds,    // Set of sockets for read events.
            fd_set* writefds,   // Set of sockets for write events.
            fd_set* exceptfds,  // Set of sockets for exceptions/errors.
            const struct timeval* timeout // Max time to block.
);
```

- Data Structures:

```
typedef struct fd_set {
    u_int fd_count;           // Number of sockets in the set.
    SOCKET fd_array[FD_SETSIZE]; // Array of sockets that are in the set.
} fd_set;

typedef struct timeval {
    long tv_sec;    // seconds
    long tv_usec;   // microseconds
} timeval;
```

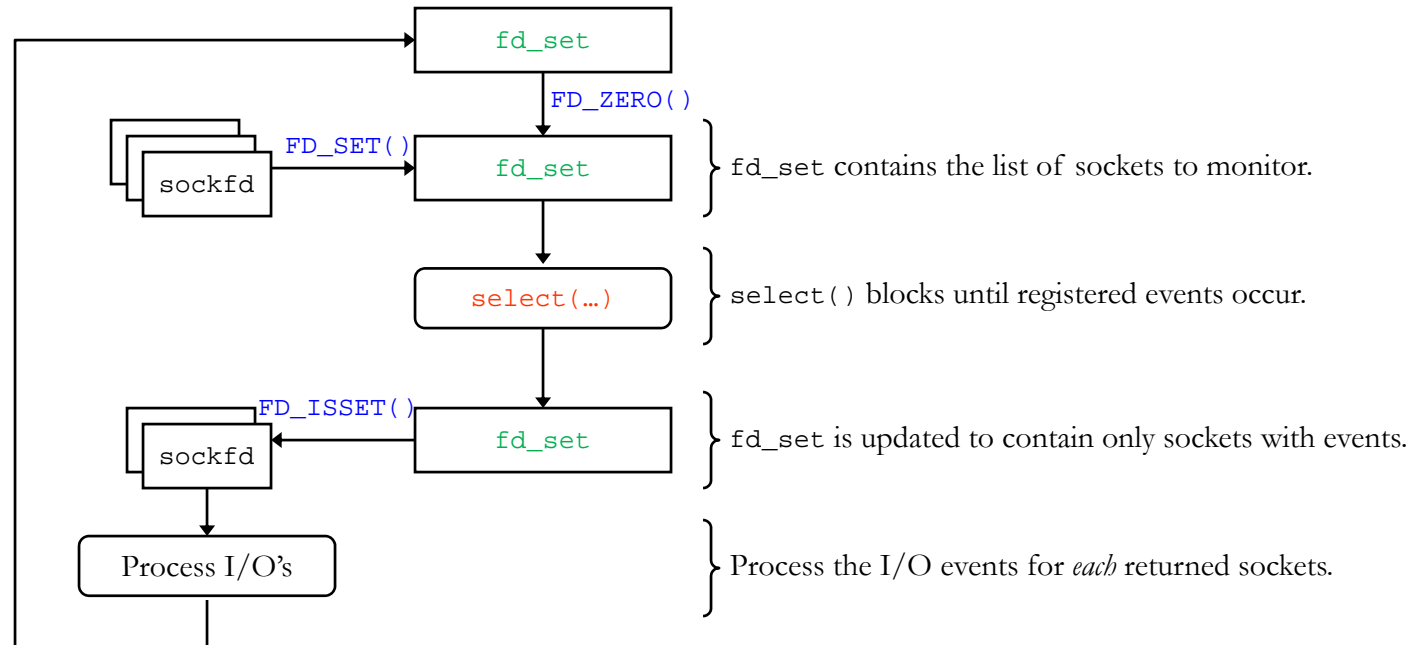
- Macros:

```
FD_CLR  (s, *set); // Removes the socket handle s from set.
FD_ISSET(s, *set); // Nonzero if s is a member of the set. Otherwise, zero.
FD_SET  (s, *set); // Adds socket handle s to set.
FD_ZERO (*set);    // Initializes the set to the null set.
```

6. I/O Multiplexing Using `select()`

Copyright Jack Y. B. Lee
All Rights Reserved

- Program Flow:



```
int select( int nfds,           // Max fd+1 in Unix/Linux, ignored in windows.
            fd_set* readfds,    // Set of sockets for read events.
            fd_set* writefds,   // Set of sockets for write events.
            fd_set* exceptfds,  // Set of sockets for exceptions/errors.
            const struct timeval* timeout // Max time to block.
);
```

7. Concurrent NetConsole

Copyright Jack Y. B. Lee
All Rights Reserved

- Implements Listener using `select ()`:

```
void ConcurrentListenerUsingSelect (char *address, char *port)
{
    /// Step 1: Prepare address structures. ///
    sockaddr_in *TCP_Addr = new sockaddr_in;
    memset (TCP_Addr, 0, sizeof(struct sockaddr_in));
    TCP_Addr->sin_family = AF_INET;
    TCP_Addr->sin_port = htons(atoi(port));
    TCP_Addr->sin_addr.s_addr = INADDR_ANY;

    /// Step 2: Create a socket for incoming connections. ///
    SOCKET Sockfd = socket(AF_INET, SOCK_STREAM, 0);
    bind(Sockfd, (struct sockaddr *)TCP_Addr, sizeof(struct sockaddr_in));
    listen(Sockfd, 5);

    /// Step 3: Setup the data structures for multiple connections. ///
    const int maxSockets = 10;          // At most 10 concurrent clients
    SOCKET socketHandles[maxSockets]; // Array for the socket handles
    bool socketValid[maxSockets];      // Bitmask to manage the array
    int numActiveSockets = 1;
    for (int i=1; i<maxSockets; i++) socketValid[i] = false;

    socketHandles[0] = Sockfd;
    socketValid[0] = true;
}
```

—continue ...—

7. Concurrent NetConsole

Copyright Jack Y. B. Lee
All Rights Reserved

- Implements Listener using `select()`:

```
/// Step 4: Setup the select() function call for I/O multiplexing. ///

fd_set fdReadSet;
while (1) {

    // Setup the fd_set //
    int topActiveSocket = 0;
    FD_ZERO(&fdReadSet);
    for (int i=0; i<maxSockets; i++) {
        if (socketValid[i] == true) {
            FD_SET(socketHandles[i], &fdReadSet);
            if (socketHandles[i] > topActiveSocket)
                topActiveSocket = socketHandles[i]; // for Linux compatibility
        }
    }


    // Block on select() //

    int ret;
    if ((ret = select(topActiveSocket+1, &fdReadSet,
                     NULL, NULL, NULL)) == SOCKET_ERROR) {
        printf("\nselect() failed. Error code: %i\n", WSAGetLastError());
        return; // continue ...
    }
}
```

7. Concurrent NetConsole

Copyright Jack Y. B. Lee
All Rights Reserved

- Implements Listener using `select ()`:

```
// Process the active sockets //
for (i=0; i<maxSockets; i++) {
    if (!socketValid[i]) continue; // Only check for valid sockets.
    if (FD_ISSET(socketHandles[i], &fdReadSet)) { // Is socket i active?
        if (i==0) { // the socket for accept() 
            // accept new connection //
            SOCKET newsfd = accept(Sockfd, 0, 0);
            // Find a free entry in the socketHandles[] //
            int j = 1;
            for ( ; j<maxSockets; j++) {
                if (socketValid[j] == false) {
                    socketValid[j] = true;
                    socketHandles[j] = newsfd;
                    ++numActiveSockets;
                    if (numActiveSockets == maxSockets) {
                        // Ignore new accept()
                        socketValid[0] = false;
                    }
                }
                break;
            }
        }
    }
}
```

continue ...

7. Concurrent NetConsole

Copyright Jack Y. B. Lee
All Rights Reserved

- Implements Listener using `select ()`:

```
        else { // sockets for recv()
            // Receive data and display to stdout. //
            char buf[256]; int len=255;
            int numread = recv(socketHandles[i], buf, len, 0);
            if (numread == SOCKET_ERROR) {
                printf("\nrecv() failed. Error code: %i\n",
                    WSAGetLastError());
                return;
            } else if (numread == 0) { // connection closed
                // Update the socket array
                socketValid[i] = false;
                --numActiveSockets;
                if (numActiveSockets == (maxSockets-1))
                    socketValid[0] = true;
                closesocket(socketHandles[i]);
            } else {
                printf("[%i]: ", i);
                for (int x=0; x<numread; x++) putchar(buf[x]);
            }
        }
        if (--ret == 0) break; // All active sockets processed.
    }
}
```

8. Final Remarks on `select ()`

Copyright Jack Y. B. Lee
All Rights Reserved

- Compatibility
 - ◆ The `select ()`-based I/O model originates from BSD sockets.
 - ◆ Does **not** require multiple processes or multi-threading
 - ◆ The Winsock implementation is slightly different but compatible.
 - ◆ Winsock also implemented another variant of `select ()`
 - `WSAAsyncSelect ()`
- Limitations
 - ◆ The maximum number of sockets that can be handled by each `select ()` is **implementation dependent**.
 - ◆ Linux can handle up to 1024 handles (minus `stdin`, `stdout`, and `stderr`).
 - ◆ Winsock varies from 1024 to very large number (e.g., 500K+ in Windows 10 Enterprise)
 - ◆ Overheads **increase** with number of sockets in `fd_set`.

8. Final Remarks on select ()

Copyright Jack Y. B. Lee
All Rights Reserved

- Setting up the fd_set in every iteration:

```
/// Step 4: Setup the select() function call for I/O multiplexing. ///
fd_set fdReadSet;
while (1) {
    FD_ZERO(&fdReadSet); int topActiveSocket = 0;
    for (int i=0; i<maxSockets; i++) {
        if (socketValid[i] == true) {
            FD_SET(socketHandles[i], &fdReadSet);
            if (socketHandles[i] > topActiveSocket)
                topActiveSocket = socketHandles[i]; // for Linux
        }
    }
}
```

- Finding out which socket has event for processing:

```
// Process the active sockets //
for (i=0; i<maxSockets; i++) { // Scan through all valid sockets
    if (!socketValid[i]) continue; // Only check for valid sockets.
    if (FD_ISSET(socketHandles[i], &fdReadSet)) { // Is socket i active?
        if (i==0) { // the socket for accept()
            // accept new connection //
            // ...(omitted)
        }
    }
}
```

8. Final Remarks on select ()

Copyright Jack Y. B. Lee
All Rights Reserved

- Benchmarking (NetConsole)

```
// Setup the fd_set //
int t=0;
timer.Start();
FD_ZERO(&fdReadSet);
t = timer.ElapseduSec();
printf("\nFD_ZERO: %i micro-seconds.\n", t);


timer.Start();
for (int i=0; i<maxSockets; i++) {
    if (socketValid[i] == true) {
        FD_SET(socketHandles[i], &fdReadSet);
    }
}
t = timer.ElapseduSec();
printf("\nFD_SET: %i micro-seconds.\n", t);

// ...
timer.Start();
if (FD_ISSET(socketHandles[i], &fdReadSet)) { // Is socket i active?
    t = timer.ElapseduSec();
    printf("\nFD_ISSET (Active): %i micro-seconds.\n", t);
}
```

8. Final Remarks on `select ()`

Copyright Jack Y. B. Lee
All Rights Reserved

- Benchmarking (NetConsole)
 - ◆ Results (time in micro-seconds)
 - FD_SET
 - ◆ 64 sockets ~ 9 micro-secs
 - ◆ 128 sockets ~ 40 micro-secs
 - ◆ 256 sockets ~ 130 micro-secs
 - ◆ 512 sockets ~ 370 micro-secs
 - ◆ 1024 sockets ~ 1400 micro-secs



Transfer Rate = x/T

Example: 512 sockets, $x=64$ bytes, $T=370\mu\text{s}$, then transfer rate = $\sim 1.4\text{Mbps}$.

References

Copyright Jack Y. B. Lee
All Rights Reserved

- Winsock2 Reference, MSDN,
<http://msdn2.microsoft.com/en-us/library/ms741416.aspx>